

DLXview

(Preliminary) User's Manual

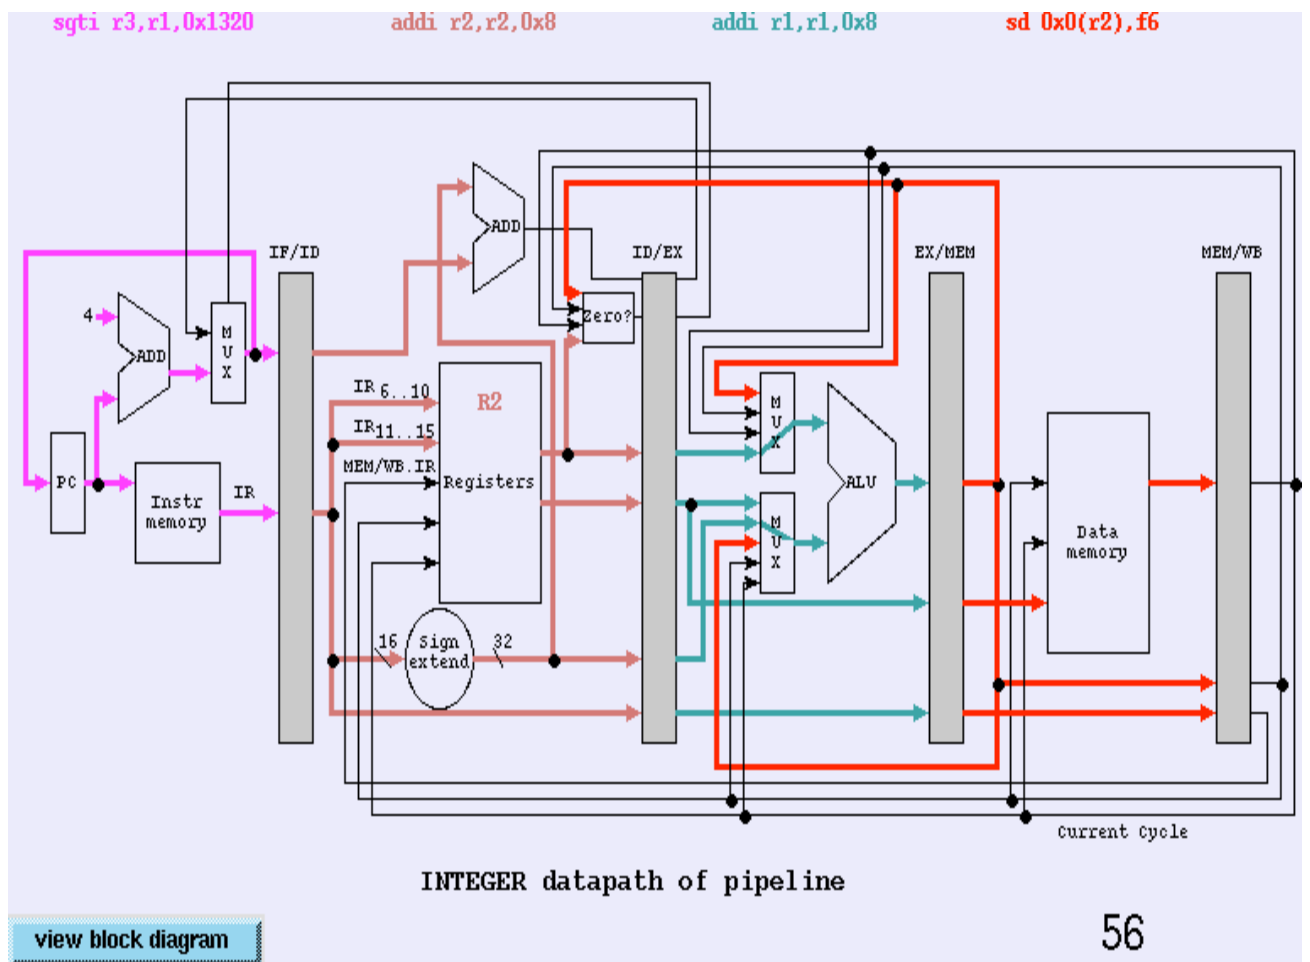


Table of Contents

| | |
|--|--|
| Introduction | |
| DLXview History and Future | |
| Contact Us | |
| Getting Started | |
| ...the Simple Way | |
| ...the Advanced Way | |
| Control Panel | |
| Configure | |
| Step Forward | |
| Next Cycle | |
| Go | |
| Step Back | |
| Previous Cycle | |
| Reset | |
| Trace | |
| Quit | |
| Pipeline Modes | |
| Basic Pipeline | |
| Tomosulo Algorithm | |
| Scoreboarding | |
| Appendix | |
| Format of a Startup File | |
| A Sample Startup File | |
| Formats of Code and Data Files | |
| Format of Register Initialization File | |
| Sample Input Files | |

Introduction

DLXview is an interactive pipelined computer simulator using the DLX instruction set described in *Computer Architecture: A Quantitative Approach (CA:AQA)* by John Hennessy and David Patterson. The main goal of dlxview is to provide a graphical environment where the internal operations of a processor are much easier to understand than their text description only. In addition to its pedagogical purpose, dlxview also serves to be a handy tool in DLX instruction set understanding, debugging, and in processor performance evaluation.

DLXview History and Future

DLXview was created by modifying and extending DLXsim, a generic, non-graphical DLX pipeline simulator, which was in turn modified from a MIPS simulator. This is not surprising given the similarities between the DLX and MIPS instruction sets. The simple pipelining model in DLXsim has been greatly enhanced to support the execution models of basic pipelining, scoreboarding, and the Tomasulo algorithm. The graphical interface built on top of Tcl/Tk tries to expose as much nitty-gritty detail as possible as explained by the CA:AQA textbook on a cycle-by-cycle basis.

The current version of dlxview is a branch of the on-going CASLE (Compiler/Architecture Simulation for Learning and Experimenting) project at Purdue University's School of Electrical and Computer Engineering, which aims to provide a comprehensive pedagogical tool for learning architectural concepts, compiler technologies, and the interactions between them. Based on the experience in the course of developing dlxview, we plan to build a more powerful simulator which will cover the general microparallel execution models such as superscalar and VLIW with differing configurations. This future tool will also target some popular commercial RISC processors, such as the Alpha and PowerPC series.

Contact Us

DLXview is a project in progress. The latest news about dlxview may be found at the dlxview (<http://yara.ecn.purdue.edu/~teamaaa/dlxview>) homepage. For project suggestions, bug reports and further information, please contact George Adams at gba@ecn.purdue.edu.

Getting Started

...the Simple Way

Be sure the "dlxview" executable is in your path, and simply type

```
dlxview
```

followed by the ENTER or RETURN key to invoke dlxview to bring up the initial dlxview simulator window.

...the Advanced Way

Additional startup files can be used to save the time of a series of interactive configuration steps. Suppose the startup file is called "foo"; there are two ways of having dlxview read it.

1. Invoke dlxview by passing the startup file as a command line option:

```
dlxview -f foo
```

followed by the ENTER or RETURN key. The disadvantage of this approach is that the input terminal will no longer accept other key strokes during the simulation session. Fortunately, this feature is irrelevant to most users who do not require extensive debugging of DLX assembly programs and dlxview internals.

2. Start a dlxview session first by typing

```
dlxview
```

followed by the ENTER or RETURN key at the command line, and type

```
source foo
```

followed by the ENTER or RETURN key after the (dlxview) prompt in the command line window.

Control Panel

In dlxview, if a button is enabled, the button text is black; if it is disabled, the button text is gray. This section explains the basic operations of the control panel window (Fig 1) buttons.

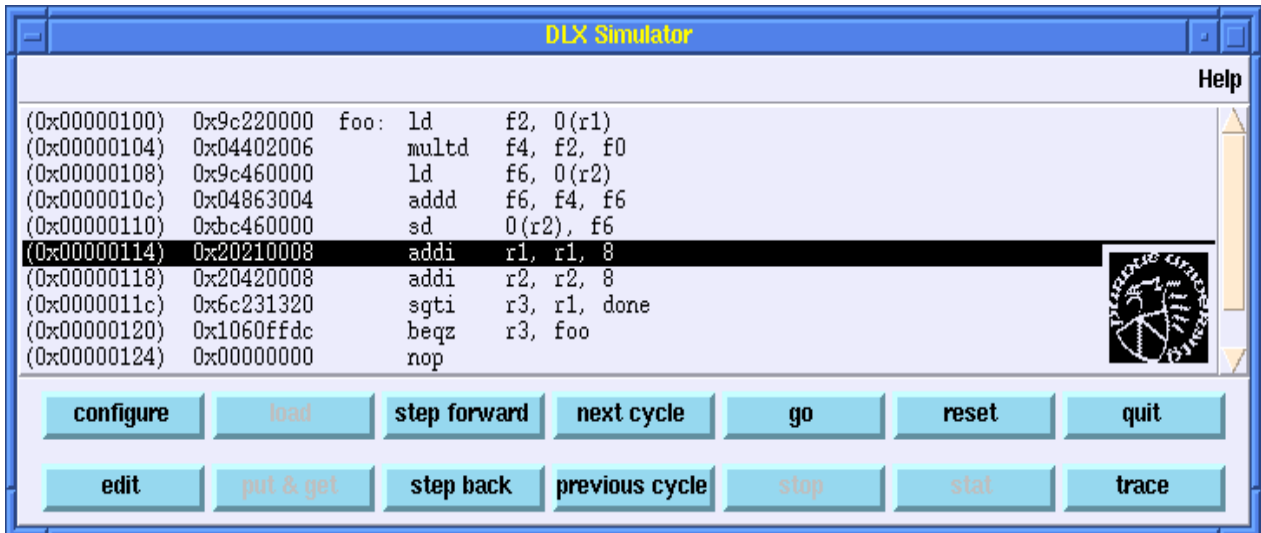


Fig 1

Configure

When dlxview is invoked without using a startup file, or right after the processor state is reset, this button must be pressed to configure the machine to be simulated.

In the mode configuration window (Fig 2), there are three radio buttons on top. Click on the button corresponding to the mode to be simulated and another larger parameter configuration window will appear corresponding to either the basic pipeline (Fig 3), scoreboard (Fig 4), or Tomasulo (Fig 5) scheduling modes. Configure the machine by dragging the sliders and clicking the radio buttons. Then exit the configuration windows by pressing the "Ok" button in the parameter configuration window, and then pressing the "Ok" button in the mode configuration window.

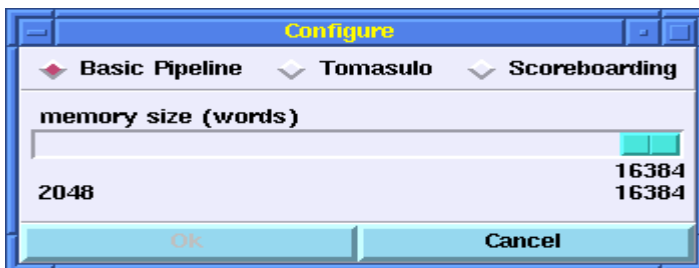


Fig 2

Basic Pipeline Configuration

| | | |
|---|--|--|
| FP adder number <input type="text" value="1"/> 1 8 | FP add stages (cycles) <input type="text" value="2"/> 1 2 5 | <input checked="" type="checkbox"/> fully piped <input type="checkbox"/> not piped |
| FP multiplier number <input type="text" value="1"/> 1 8 | FP mul stages (cycles) <input type="text" value="5"/> 1 5 10 | <input checked="" type="checkbox"/> fully piped <input type="checkbox"/> not piped |
| FP divider number <input type="text" value="2"/> 0 2 8 | FP div latency (cycles) <input type="text" value="19"/> 1 19 50 | <input type="checkbox"/> fully piped <input checked="" type="checkbox"/> not piped |
| <input type="button" value="Ok"/> <input type="button" value="Cancel"/> | | |

Fig 3

Scoreboarding Configuration

| | |
|--|---|
| integer unit number <input type="text" value="1"/> 1 8 | integer unit latency (cycles) <input type="text" value="1"/> 1 3 |
| FP add reservation station number <input type="text" value="2"/> 1 2 8 | FP add latency (cycles) <input type="text" value="2"/> 1 2 5 |
| FP multiply reservation station number <input type="text" value="2"/> 1 2 8 | FP multiply latency (cycles) <input type="text" value="5"/> 1 5 10 |
| FP divide reservation station number <input type="text" value="2"/> 0 2 8 | FP divide latency (cycles) <input type="text" value="19"/> 1 19 50 |
| <input type="button" value="Ok"/> <input type="button" value="Cancel"/> | |

Fig 4

The image shows a 'Tomasulo Algorithm Configuration' window with the following settings:

| integer reservation station number | integer op latency (cycles) |
|--|-------------------------------|
| 1 2 | 1 3 |
| 8 | |
| FP add reservation station number | FP add latency (cycles) |
| 1 1 | 1 2 |
| 8 | 5 |
| FP multiply reservation station number | FP multiply latency (cycles) |
| 1 1 | 1 5 |
| 8 | 10 |
| FP divide reservation station number | FP divide latency (cycles) |
| 0 2 | 1 19 |
| 8 | 50 |
| load buffer number | load buffer latency (cycles) |
| 0 0 | 1 1 |
| 8 | 20 |
| store buffer number | store buffer latency (cycles) |
| 0 0 | 1 1 |
| 8 | 20 |

At the bottom are 'Ok' and 'Cancel' buttons.

Fig 5

In the middle of a simulation session, the "configure" button in the control panel can be pressed again to check the current configurations. However, the sliders and buttons on the configuration windows are disabled; i.e., the processor configuration cannot be changed once simulation has started. To abort the current configuration, press the "reset" button on the control panel first, and then configure the machine from there.

Note that in the CA:AQA text, the latency of a functional unit is defined as the number of intervening cycles between an instruction that produces a result and an instruction that uses the result (see CA:AQA, page 189). For example, the integer unit will have latency 0, which is somewhat unintuitive. In dlxview, the latency of a functional unit is defined as the total number of clock cycles an instruction will stay in this functional unit. Therefore, the latency is one clock cycle more than specified in CA:AQA. For more information on the configurable parameters, please refer to tables in the Appendix.

The processor must be configured before loading a file. A file selection window will pop up after pressing the "load" button. DLXview accepts three types of files before a simulation session begins. The *.s files are assembly code file, *.d files are

data file. It is not necessary to have a separate data file. Code and data files can be merged into a single *.s file. *.i files contain commands to initialize the registers to be used. It is possible to add a few instructions in the assembly code to do the same job, the catch is that the code segment won't start at clock cycle one. Register initialization is unavoidable in some situations, such as examining the execution of a loop in which some instructions have pre-specified operands. Letting this loop start at clock cycle one, as most of the examples and exercises in the CA:AQA text do, will make it easier to focus on the loop instructions themselves without worrying about other minor details such as the dependence between a loop instruction in the first iteration of a loop and an instruction outside the loop. It is much handier to write a separate register initialization file for these cases.

When loading related input files, clicking the left mouse button on the file selection window for any one of the related files will highlight these files as a group, if the related files have the same base name. Press the "load" button in the "Load" window to load these files. When choosing to load a single file, the contents of the selected file will appear to the right of the file selection window. You can load as many files as you wish and switch between related and single load modes. When all the files are loaded, press the "Done" button to close the "Load" window. Note that pressing "Done" directly without pressing "Load" first will result in no file being loaded.

The "edit" button can be invoked any time during a simulation session to revise the current DLX code and edit the input files. When the edit window pops up, the assembly code executed by the current simulation session will be automatically copied into the edit window. If a totally new assembly file is going to be edited, simply press the "Clean" button in the edit window to get rid of the old files. When editing is completed, the new file can be directly loaded into dlxview to start another simulation session by pressing the "Load" button in the edit window. Before the new file is actually loaded, a dialog box will ask if any other data file needs to be loaded first. If the answer is "No", then the *.d file used in the current simulation session will be loaded first by default. Once a newly edited file has been loaded into the simulator through the edit window, the current simulation session is interrupted, and the processor is reset in order to run the latest loaded file. With the help of the code editor, you can change the assembly code being executed at any time and try the new simulation without leaving the dlxview environment.

Step Forward

If no instruction has entered the pipeline yet, pressing "step forward" will pop up a window to allow entering the starting address. Otherwise, the simulation will continue from wherever execution previously stopped and execute the next instruction. The clock cycle will always advance to the first stage of the newly processed instruction.

Next Cycle

If no instruction has entered the pipeline yet, pressing "next cycle" will pop up a window to allow entering the starting address. Otherwise, the simulation will advance by one clock cycle. The simulator will not necessarily proceed to the next instruction, because of pipeline stalls. Pressing "next cycle" continuously will show every detail in the simulated pipeline.

Go

If no instruction has entered the pipeline yet, pressing "go" will pop up a window to allow entering the starting address. Otherwise, the simulation will continue from wherever execution previously stopped and will not stop until the execution terminates naturally. Do not forget to put a trap #0 instruction at the end of the assembly code; otherwise, the simulator will continuously fetch new nop instructions from the zeroed memory locations preceding the input code.

Step Back

This button provides the flexibility to back step one instruction and view the state of the processor at the clock cycle when the last instruction entered the pipeline. If more than 1K clock cycles have already been simulated, it will take a little while to move the pipeline to an earlier state. Be patient!

Previous Cycle

The button provides the flexibility to wrap back the state of the processor one clock cycle. If more than 1K clock cycles have already been simulated, it will take a little while to move the pipeline to an earlier state. Be patient!

Reset

The "reset" button allows changing the pipelining mode, the configuration of the processor, or the input code.

Trace

The "trace" button provides the capability to output an instruction trace or memory reference trace. This may be fed into a trace-driven simulator or the dinero cache simulator later to do further analysis. Enter the file name to save the trace as. The trace will be collected starting from the current instruction. Normally, the trace file is not available until the current simulation session ends; i.e., the simulation stops naturally or is interrupted by a reset. However, when pressing either the "step back" or "previous cycle" buttons during the simulation, the tracing process is stopped automatically to avoid producing garbage wrapped around trace. In this case, a trace file containing the information from where tracing started to the last instruction will be generated.

Quit

Terminates the dlxview session.

Pipeline Modes

Basic Pipeline

In the visualization window (Fig 6) for the basic pipeline, the top half is an instruction and clock cycle table, and the bottom half can be switched between the block diagram of the DLX pipeline (see Figure 3.44 in the CA:AQA text) and the datapath of the integer unit (see Figure 3.22 in the CA:AQA text). Each instruction entering the pipeline is assigned a unique color. Instructions of a particular color listed in the top half of the window execute in the pipeline stage corresponding to same color. The window exhibits the pipeline state for one particular clock cycle. This clock cycle number is displayed at the lower-right corner of the visualization window, and is highlighted by a gray strip in the clock cycle table.

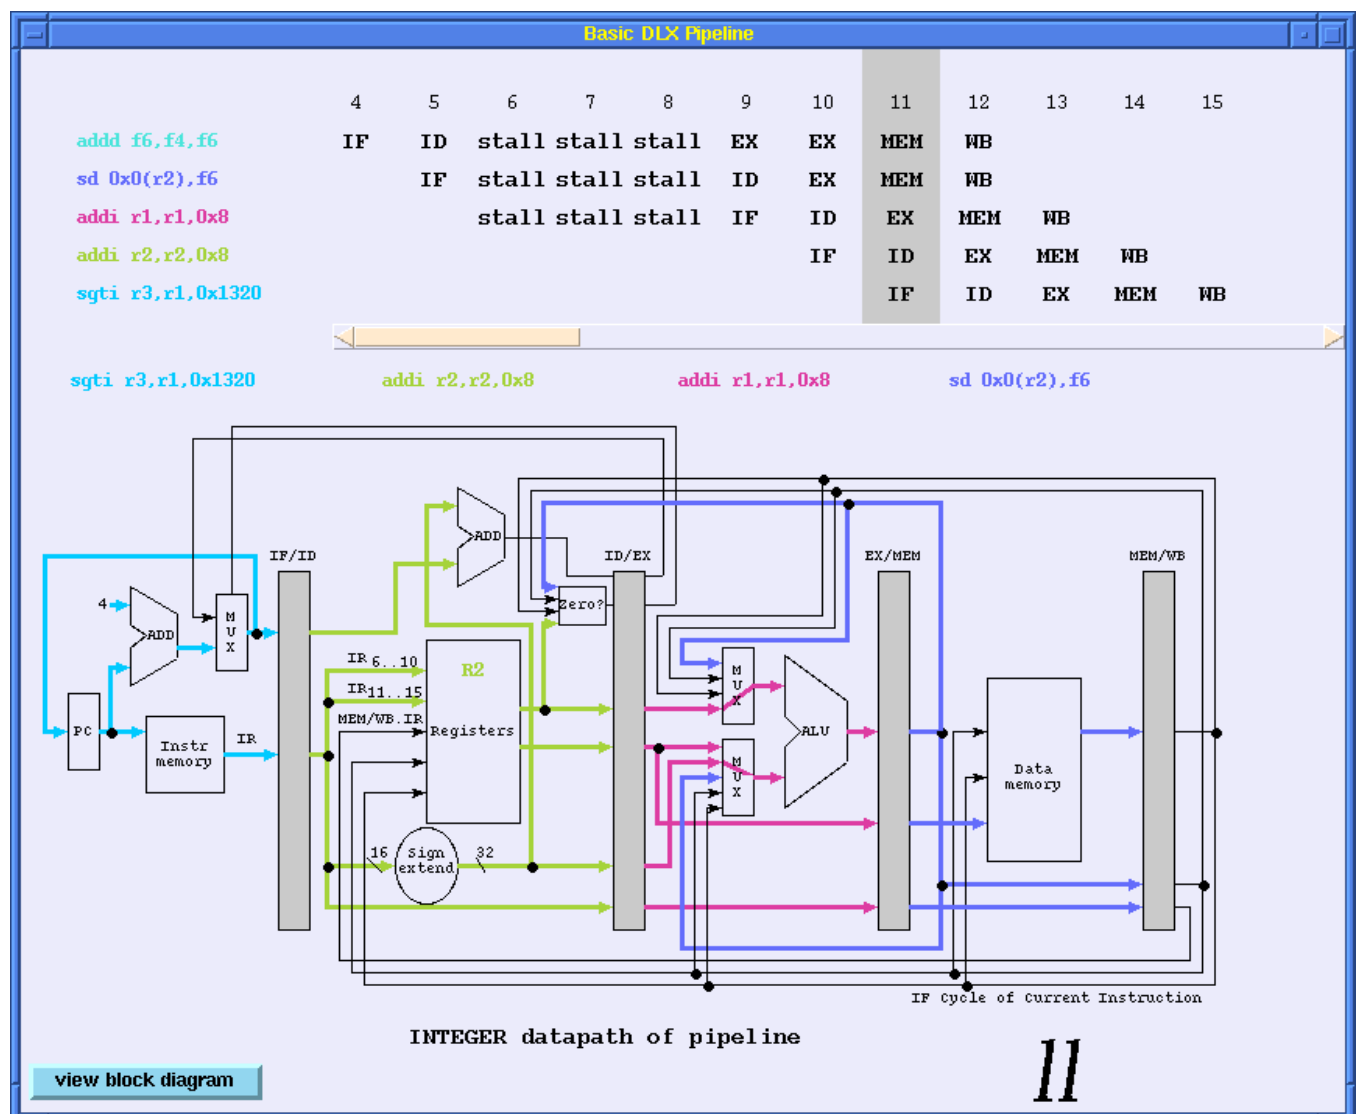


Fig 6

The instruction and clock cycle table are capable of displaying the latest five instructions simulated by dlxview, with the last row showing the last instruction entering the pipeline. The highlighted gray strip identifies which instruction is in which pipeline stage. There might be more than five instructions in the pipeline at any time; the earlier instructions are wrapped off in this table. To determine exactly how many instructions are in the pipeline, count the number of colored boxes in the bottom pipeline visualization.

In the pipeline block diagram, if a floating point functional unit is fully pipelined, the EX pipe stages for this functional unit will be drawn as separated boxes. On the other hand, the multicycle EX stage of a unpipelined floating point unit is shown as a long box with internal lines. The block diagram as a whole is also convenient for determining which instruction is in which pipeline stage. For a pipeline that supports multiple outstanding FP operations, there might be multiple instructions entering the MEM or WB stage simultaneously. (It can be proved that for the particular DLX machine described in the section 3.7 of the CA:AQA text, at most two instructions will go to the MEM or WB stages at the same time. This is left as an exercise for the reader.) In this case, multiple colors will occupy the MEM and WB boxes (they should be split in a real implementation).

The integer datapath shows what is happening in the integer datapath. Nevertheless, floating point operations are exhibited here also when they are in the IF and ID stages. At a particular clock cycle, there might be some floating point instruction executing or writing their results back to the register file; they may not appear in the integer datapath.

When a mux is used for selecting operands, the line connecting the output and the selected input indicates where the operand came from. For example, by examining the input muxes to the ALU, you can determine whether operands came from the register file, the forwarding path, or the immediate operand stored in the ID/EX stage-register. If an operand is forwarded, it is also easy to trace back to the forwarding source from the highlighted wires.

The following is a brief description of the features of the DLX pipeline defined in the CA:AQA text that are supported by dlxview.

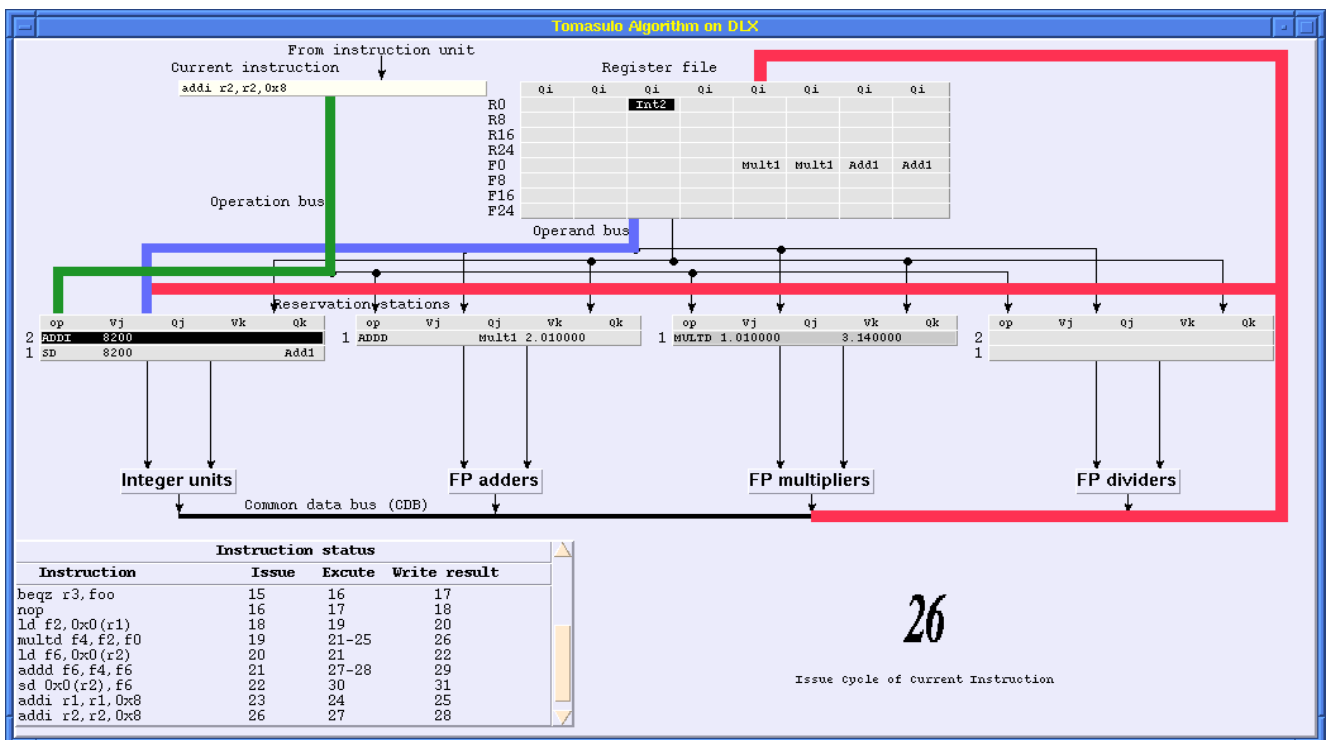
1. Instructions are fetched and decoded sequentially.
2. Multiple floating point functional units, either fully pipelined or not, can be configured for the processor. However, only one integer unit is allowed.
3. Branches are resolved at the ID stage. There is one branch delay slot after the branch instruction, and branch penalties are reduced by the delayed-branch scheme. It is up to the compiler or user, whichever provides the DLX code, to make sure the successor instruction of a branch is valid and useful. If no useful instructions can be found, a nop instruction should be inserted after the branch in the assembly code.
4. There is only one write port to the FP register file.
5. Four types of hazards are checked for in the ID stage:

1. Structural hazards,
2. WAR hazards,
3. WAW hazards, and
4. FP register file write port conflicts.

Moving branch handling logic from the MEM stage to the ID stage may cause additional stalls in the pipeline because of data dependences. Since this hazard is not dlxviewussed in the CA:AQA text, the dlxview implementation is briefly explained here. The hazard may be viewed by writing a proper code sequence and testing it with dlxview.

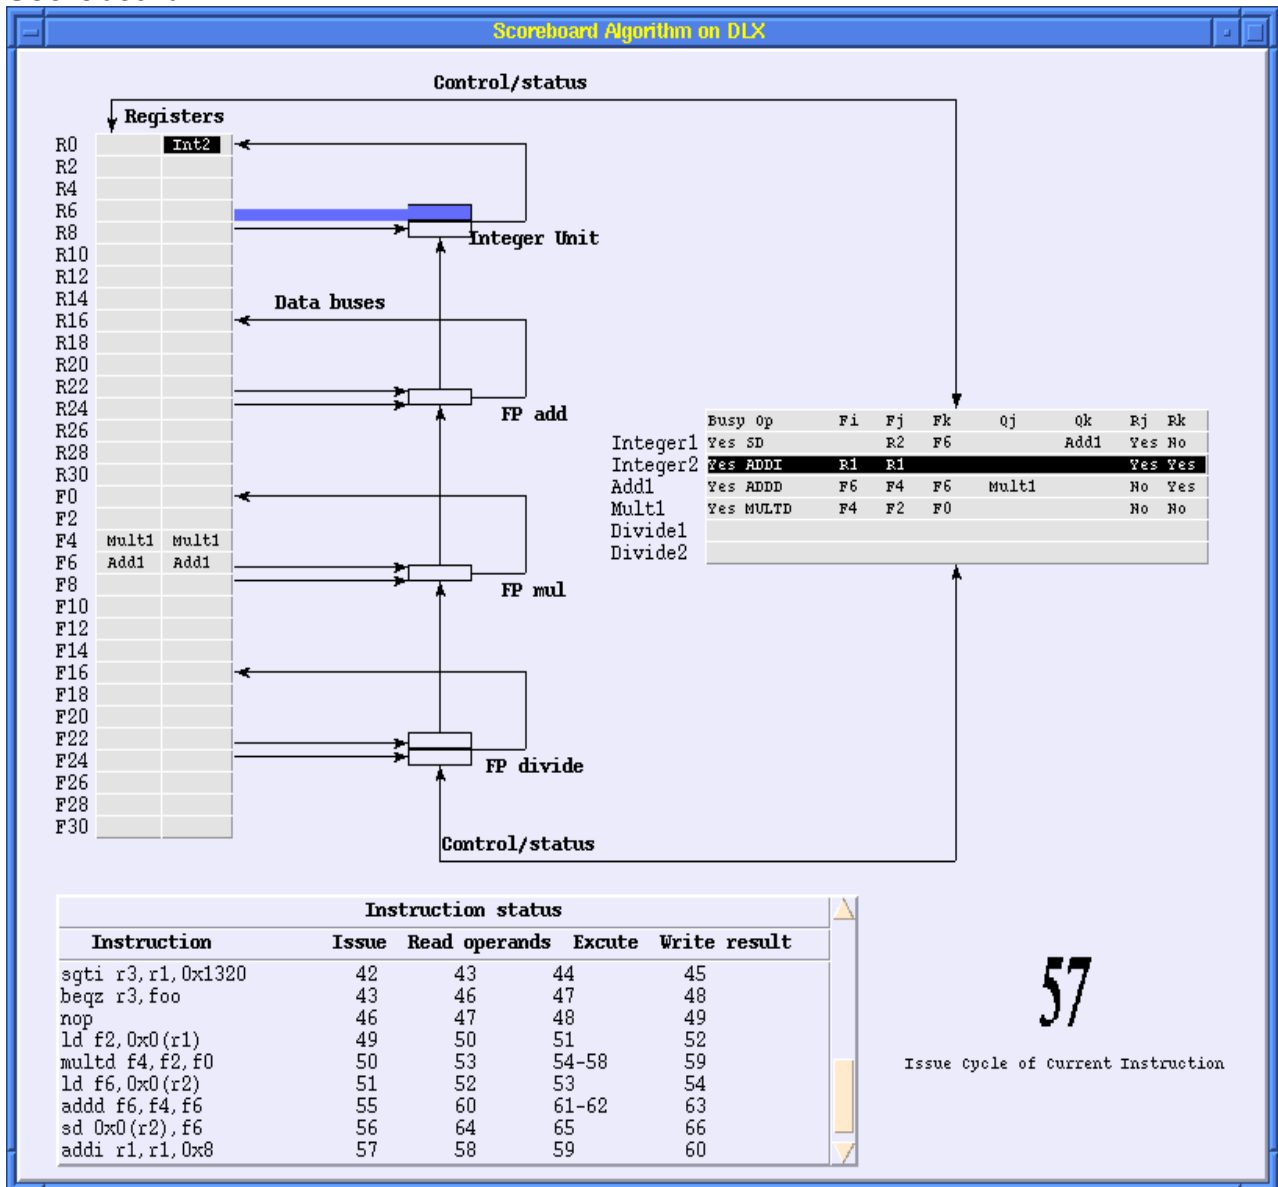
1. If the operand of a branch is generated by the preceding integer ALU instruction, the pipeline will stall for one cycle. After that, the operand is forwarded from the EX/MEM stage. The stall might last longer for FP branches in this case, depending on the latency of FP instruction.
2. If the operand of a branch is generated by the preceding load instruction, the pipeline will stall for two cycles. After that, the operand is forwarded from the MEM/WB stage.
3. If the operand of a branch is generated by the load instruction preceding the instruction which precedes the branch instruction, the pipeline will stall for one cycle. After that, the operand is forwarded from the MEM/WB stage-register.

Tomasulo Algorithm



visualization window

Scoreboard



visualization window

Appendix

Format of a Startup File

The startup file contains a sequence of parameter assignments. Each assignment is a variable-value pair with the following Tcl syntax:

```
varName ?value?
```

If there is more than one item in a value (for example, multiple assembly code files and data files are required to be loaded before execution) then these items should be wrapped by a pair of double quotes or curly parenthesis. There is no ordering of assignments; however, the end of assignments should be followed by a

```
Start
```

command. The configurable variables are listed in Table 1. All the variables and commands are case sensitive: make sure to type them correctly. Some variables are only meaningful in certain pipelining modes; Table 1 gives the applicable mode to each variable in the last column. Unpredictable results or core dumps will appear if a variable is assigned under a wrong mode. A "must" means the corresponding variable must be set no matter what pipelining mode has been chosen. If a variable is not been specified in the startup file, it will take the default value. Table 2 gives the domain of each variable and their default values.

Table 1. Configurable Variables and Their Meanings

| VARIABLE | MEANING | PIPELINE MODE |
|----------------|--------------------------------|-------------------------------|
| pipeline_mode | pipelining mode | must |
| input_files | code/data/initialize file name | must |
| start_address | start address in the code | BasicPipe,ScoreBoard,Tomasulo |
| numWords | memory size | BasicPipe,ScoreBoard,Tomasulo |
| num_int_units | number of integer units | ScoreBoard,Tomasulo |
| num_add_units | number of FP adders | BasicPipe,ScoreBoard,Tomasulo |
| num_mul_units | number of FP multipliers | BasicPipe,ScoreBoard,Tomasulo |
| num_div_units | number of FP dividers | BasicPipe,ScoreBoard,Tomasulo |
| int_latency | integer unit latency | BasicPipe,ScoreBoard,Tomasulo |
| fp_add_latency | FP adder latency | BasicPipe,ScoreBoard,Tomasulo |
| fp_mul_latency | FP multiplier latency | BasicPipe,ScoreBoard,Tomasulo |
| fp_div_latency | FP divider latency | BasicPipe,ScoreBoard,Tomasulo |
| add_ful_pipe | FP adder fully pipelined? | BasicPipe |
| mul_ful_pipe | FP multiplier fully pipelined? | BasicPipe |
| div_ful_pipe | FP divider fully pipelined? | BasicPipe |
| num_load_bufs | number of load buffers | Tomasulo |

| | | |
|-------------------|-------------------------|----------|
| num_store_bufs | number of store buffers | Tomasulo |
| load_buf_latency | load buffer latency | Tomasulo |
| store_buf_latency | store buffer latency | Tomasulo |

Table 2. Domain and Default Values for Configurable Variables

| VARIABLE | DOMAIN | DEFAULT |
|-------------------|-------------------------------|---------|
| pipeline_mode | BasicPipe,ScoreBoard,Tomasulo | none |
| input_files | files with .s .d .i extension | none |
| start_address | any legal address or symbol | 0x1000 |
| numWords | [1, 16384] | 16384 |
| num_int_units | 1 | 1 |
| num_add_units | [1, 8] | 2 |
| num_mul_units | [1, 8] | 2 |
| num_div_units | [0, 8] | 2 |
| int_latency | [1,) | 1 |
| fp_add_latency | [1,) | 2 |
| fp_mul_latency | [1,) | 5 |
| fp_div_latency | [1,) | 19 |
| add_ful_pipe | 0 for no, 1 for yes | 1 |
| mul_ful_pipe | 0 for no, 1 for yes | 1 |
| div_ful_pipe | 0 for no, 1 for yes | 0 |
| num_load_bufs | [0, 8] | 0 |
| num_store_bufs | [0, 8] | 0 |
| load_buf_latency | [1,) | 1 |
| store_buf_latency | [1,) | 1 |

A Sample Startup File

```

set pipeline_mode BasicPipe
set fp_add_latency 3
set num_add_units 1
set mul_ful_pipe 0
set fp_mul_latency 10
set num_div_units 1
set num_mul_units 1
set fp_div_latency 24
set input_files {f347.i f347.d f347.s}

```

Formats of Code and Data Files

1. Comments are started with a semicolon, and continue to the end of line.
2. Constants can be entered either with or without a preceding number sign.
3. The format of instructions and their operands are as shown in the CA:AQA text.

4. Labels are defined by a group of non-blank characters starting with either a letter, an underscore, or a dollar sign, and followed immediately by a colon. They are associated with the address immediately following the last block of information stored. This has the negative effect that if code follows a label following a block of data that does not end on a word boundary (multiple of 4), the label will not point to the first instruction in the code, but instead to 1 to 3 bytes before (since the address is only rounded when it is necessary to correctly align data). This is done so that if a label is found in the middle of a data section, it will point to the start of the next section of data without the data having to be aligned to a word boundary. The work-around for this is to use the `.align` (see below) directive before labels that will not be aligned with the data following them. Labels can be accessed anywhere else within that file, and in files loaded after that if the label is declared as `.global` (see below).

5. The text (code) is ended by a `trap #0` instruction.

While the assembler is processing an assembly file, the data and instructions it assembles are placed in memory based on either a text (code) or data pointer. Which pointer is used is selected not by the type of information, but by whether the most recent directive was `.data` or `.text`. The program initially loads into the text segment. The assembler supports several directives which affect how it loads the DLX memory. These should be entered in the place where you would normally place the instruction and its arguments. The directives currently supported by dlxview are:

- `.align n`
Cause the next higher address with the lower n bits zeroed (the next closest address greater than or equal to the current address that is a multiple of 2^n).
- `.ascii "string1", "string2", ...`
Store the strings listed on the line in memory as a list of characters. The strings are not terminated by 0 byte.
- `.asciiz string1", "string2", ...`
Similar to `.ascii`, except each string is followed by a 0 byte (like C strings).
- `.byte "byte1", "byte1", ...`
Store the bytes listed on the line sequentially in memory.
- `.data [address]`
Cause the following code and data to be stored in the data area. If an address was supplied, the data will be loaded starting at that address, otherwise, the last value for the data pointer will be used. If we were just reading code based on the text (code) pointer, store that address so that we can continue from there later (on a `.text` directive).
- `.double "number1", "number2", ...`
Store the numbers listed on the line sequentially in memory as double precision floating point numbers.
- `.float "number1", "number2", ...`
Store the numbers listed on the line sequentially in memory as single precision floating point numbers.

- `.global label`
Make the label available for reference by code found in files loaded after this file.
- `.space size`
Move the current storage pointer forwarded by size bytes (to leave some empty space in memory).
- `.text [address]`
Cause the following code and data to be stored in the text (code) area. If the address was supplied, the data will be loaded starting at that address, otherwise, the last value for the text pointer will be used. If we were just reading code based on the data pointer store that address so that we can continue from there later (on a `.data` directive).
- `.word "word1", "word1", ...`
Store the words listed on the line sequentially in memory.

DLXview allows the user access to a few simple C library functions through the use of the trap operation. Currently supported functions are: **open ()** (trap #1), **close ()** (trap #2), **read ()** (trap #3), **write ()** (trap #4), **printf ()** (trap #5). When the appropriate trap is invoked, the first argument should be located in the word starting at the address in r14, with the following arguments (as seen in a C statement calling the function) in words above that (r14+4, r14+8, ...). The result from the function call will be placed in r1 (this means there is currently no support for the library functions that return floating point values). If a double precision floating point value is to be passed to a library function, it will occupy two adjacent words with the lower word the value of the even valued floating point register, and the higher word containing the value of the odd valued floating point register (f0 in 0(r14), f1 in 4(r14)).

Note: The dlxview file loader is completely inherited from DLXsim, and this help section is directly copied from the DLXsim User's Manual.

Format of Register Initialization File

Two types of register initialization commands exist. The first type is used to initialize integer registers, and the second type is used to initialize floating point registers.

For register designators, any of the names "r0" through "r31" and "f0" through "f31" may be used. The names "\$0" through "\$31" may also be used (instead of "r0" through "r31"), however, it is safer to use "r" instead of "\$" to avoid the confusion with Tcl variables. The name "pc" may be used to refer to the program counter.

- `put reg number`
Store number in the register given by reg. To store floating point numbers (single or double precision), use the `fput` command.
- `fput reg number [precision]`

Store number in the register given by reg. If precision is d, the number is stored as a double precision floating point number (in two words). If precision is f, or no precision is given, the number is stored as a single precision floating point number.

Sample Input Files

f347.s

```
LL:  ld      f2,0(r1)
     multd   f4, f2, f0
     ld      f6, 0(r2)
     addd    f6, f4, f6
     sd      0(r2), f6
     addi    r1, r1, 8
     addi    r2, r2, 8
     sgti    r3, r1, done
     beqz    r3, LL
     nop
     trap   #0
```

f347.i

```
put    r1    0x1000
put    r2    0x2000
fput   f0    3.14  d
```

f347.d

```
.data 0x1000
.double 1.00, 1.01, 1.02, 1.03, 1.04, 1.05, 1.06, 1.07, 1.08, 1.09
.double 1.10, 1.11, 1.12, 1.13, 1.14, 1.15, 1.16, 1.17, 1.18, 1.19
.double 1.20, 1.21, 1.22, 1.23, 1.24, 1.25, 1.26, 1.27, 1.28, 1.29
.double 1.30, 1.31, 1.32, 1.33, 1.34, 1.35, 1.36, 1.37, 1.38, 1.39
.double 1.40, 1.41, 1.42, 1.43, 1.44, 1.45, 1.46, 1.47, 1.48, 1.49
.double 1.50, 1.51, 1.52, 1.53, 1.54, 1.55, 1.56, 1.57, 1.58, 1.59
.double 1.60, 1.61, 1.62, 1.63, 1.64, 1.65, 1.66, 1.67, 1.68, 1.69
.double 1.70, 1.71, 1.72, 1.73, 1.74, 1.75, 1.76, 1.77, 1.78, 1.79
.double 1.80, 1.81, 1.82, 1.83, 1.84, 1.85, 1.86, 1.87, 1.88, 1.89
.double 1.90, 1.91, 1.92, 1.93, 1.94, 1.95, 1.96, 1.97, 1.98, 1.99
.global done
done:  .double 0

.data 0x2000
.double 2.00, 2.01, 2.02, 2.03, 2.04, 2.05, 2.06, 2.07, 2.08, 2.09
.double 2.10, 2.11, 2.12, 2.13, 2.14, 2.15, 2.16, 2.17, 2.18, 2.19
.double 2.20, 2.21, 2.22, 2.23, 2.24, 2.25, 2.26, 2.27, 2.28, 2.29
```

.double 2.30, 2.31, 2.32, 2.33, 2.34, 2.35, 2.36, 2.37, 2.38, 2.39
.double 2.40, 2.41, 2.42, 2.43, 2.44, 2.45, 2.46, 2.47, 2.48, 2.49
.double 2.50, 2.51, 2.52, 2.53, 2.54, 2.55, 2.56, 2.57, 2.58, 2.59
.double 2.60, 2.61, 2.62, 2.63, 2.64, 2.65, 2.66, 2.67, 2.68, 2.69
.double 2.70, 2.71, 2.72, 2.73, 2.74, 2.75, 2.76, 2.77, 2.78, 2.79
.double 2.80, 2.81, 2.82, 2.83, 2.84, 2.85, 2.86, 2.87, 2.88, 2.89
.double 2.90, 2.91, 2.92, 2.93, 2.94, 2.95, 2.96, 2.97, 2.98, 2.99